

Contents

What is <i>actually</i> in the <code>.fsf</code> extension on Software Heritage?	1
1. Motivation & questions	1
2. Method	1
3. Experiments	2
4. Results	2
5. Key findings & lessons (summary)	7
6. Toward a general methodology	7
7. Limitations	7
8. Reproducibility & artefacts	8

What is *actually* in the `.fsf` extension on Software Heritage?

An empirical study of the `.fsf` file extension in the Software Heritage (SWH) archive — the world’s largest source-code archive. Toolkit: `tools/fsf/`.

Bottom line. `.fsf` on SWH is, almost entirely, the **FSL FEAT design file** — a configuration for fMRI (functional MRI) neuroimaging analysis, written in **Tcl set** syntax. It is not a general-purpose programming language^a; it is best read as a domain-specific / embedded configuration language hosted in Tcl. Reading content also reveals *polysemy*: the same extension is occasionally a **GLSL fragment shader** or a data-management pointer — different files, different languages.

^a“Not a programming language” is a statement about the *file’s role*, not its syntax. An FSL FEAT design is a **declarative configuration** that parametrises an analysis pipeline — it does not, by itself, compute. It is, however, *expressed in Tcl* (a real programming language) as a sequence of `set fmri(...)` **value** statements, so it can equally be described as an **internal / embedded DSL over Tcl**. Throughout, we distinguish *content role* (configuration / DSL vs. general-purpose program) from *host syntax* (Tcl).

1. Motivation & questions

A file extension is the cheapest possible label. Software Heritage has archived tens of thousands of files ending in `.fsf`, but the extension alone says nothing about what they contain, who wrote them, or why. We ask:

- **Q1 — What is it?** What kind of content is `.fsf`, in what format, and which language(s)/notation(s) does it relate to?
- **Q2 — Is it a programming language, and how pure is the extension?** Is all `.fsf` one thing, or is the extension polysemous / contaminated?
- **Q3 — What is the sub-structure?** If it is a specific format, what are its variants and how were the files produced?
- **Q4 — Provenance.** Which communities and repositories does it come from?

2. Method

We use a general content-characterisation pipeline (the same one we apply to other extensions):

Stage	What it does
Sample	dedup by content <code>sha1_git</code> ; draw N contents (two sampling frames, §3)
Fetch	pull raw bytes from SWH by <code>sha1_git</code> (self-verifying), cached locally
Indicators	deterministic metrics from the bytes: line/comment counts, <code>set fmri(...)</code> directives, FEAT level, EV count, format
LLM-judge	an enum-constrained verdict via OpenRouter (Claude Sonnet 4.6, strict <code>json_schema</code>)
Reclassifier	a zero-API content rule, validated against the LLM

Stage	What it does
Origins	recover the repository each content came from
Review	a web app over every label (tools/fsf/review_app.py) for human ground truth

The key design choice is *what we ask the judge*: not “is this language X?” but “**what is this content, and which languages/notations does it relate to?**” — so it leads with `content_type`, a free-text format, `expressed_in`, a `related_languages[]` list, and `ecosystem_tool`, with an `is_programming_language` flag and format-specific detail as secondary fields. This is what lets a *polysemous* extension reveal itself.

3. Experiments

Two sampling frames plus a classifier evaluation:

#	Experiment	Sample (frame)	N judged	What it shows
E1	File-level population	1 000, uniform-random	982	what a typical <code>.fsf</code> <i>file</i> is
E2	Project-level population	757, origin-diverse (≤ 1 file per repo)	706	what a typical <code>.fsf</code> <i>project</i> is (no repo dominates)
E3	Cheap-classifier eval	reclassifier vs the LLM judge	1 662	whether a zero-API marker can replace the judge

Fixed throughout: model `anthropic/claude-sonnet-4.6`, temperature 0, structured outputs. Total **1 662 judged contents** (from a pool of 21 800 across 757 repositories), ~\$31.

4. Results

4.1 What `.fsf` is — an FSL FEAT configuration in Tcl (Q1, Q2)

A fetched file looks like this:

```
# FEAT version number
set fmri(version) 3.14
set fmri(inmelodic) 1
set fmri(level) 1
set fmri(analysis) 7
set fmri(outputdir) "..."
```

This is an **FSL FEAT design file** — the configuration format of *FEAT* (fMRI Expert Analysis Tool) in the FMRI Software Library, the dominant open-source fMRI analysis suite. It is a heavily-commented `Tcl set fmri(...)` **config**, usually generated by the FEAT GUI.

- **Content type:** `config-or-dsl` in **76–85 %** of files; **other** 13–16 %; small data / docs / source-code tails.
- **Is it a programming language?** `is_programming_language = False` for **1 641 / 1 662 = 99 %** (in the sense of the footnote above). The method correctly identifies a configuration format and does not mistake its Tcl host syntax for “a program”.
- **Ecosystem:** FSL / FEAT (FMRI), overwhelmingly.

Key finding — **what `.fsf` is**. `.fsf` on SWH is the FSL FEAT fMRI-analysis design file: a neuroimaging **configuration** (99 % not a general-purpose program) **expressed in Tcl** — an embedded DSL over Tcl.




Figure 1: Content type

4.2 What it relates to — polysemy (Q1, Q2)

The most informative output is *which languages the content relates to*.

Tcl 77 % (the FEAT host syntax), then a real tail: **XML** (57), **GLSL** (11), **C++** (6), **ML** (6), **HTML/CSS/JavaScript** (5 each), **C** (5), **OCaml/Haskell** (4). So **.fsf** *relates to* Tcl by construction, but the extension is genuinely **polysemous** across the archive.

Key finding — polysemy. **.fsf** is not one thing. Its dominant meaning is FSL FEAT (Tcl), but the same extension is *also* used for **GLSL fragment shaders** — indeed the ~21 files the judge marks as a real programming language are exactly the shaders (`2DShader.fsf`, `ppFragment.fsf`) and C++/GLSL headers. A binary “is it language X?” question would have hidden this entirely; asking “what does it relate to?” surfaces it.

4.3 FSL FEAT sub-types (Q3)

Among the FEAT files: **first-level** analyses dominate (55–69 %) over higher-level (12–16 %); **task-GLM** designs (48–62 %) over resting-state / MELODIC-ICA (7 %) and group-stats (12–16 %); and they are mostly **GUI-generated** (54–76 %) rather than hand-edited.

4.4 The contamination / polysemy tail (Q2)

The **other** slice (13–16 %) and the non-Tcl languages are where **.fsf** is *not* FSL FEAT:

- **GLSL fragment shaders** — **.fsf** is also a *fragment-shader* extension (**.fs/.fsf**); some are C++ headers embedding GLSL.
- **git-annex / DataLad pointers** — a *large* artefact: ~120 files per frame are git-annex symlink/pointer stubs, not design content. In DataLad-managed neuroimaging datasets the real **.fsf** is *annexed*, so what SWH archived at that path is a pointer. This is ecosystem-specific contamination.
- Odd tails: a fractal-rendering config, stray XML/HTML/JSON.




Figure 2: What .fsf relates to




Figure 3: FEAT sub-types

Lesson — contamination is ecosystem-specific. Every extension carries its own kind of non-target content. Here it is git-annex pointers (a data-management artefact of the neuroimaging community) and GLSL shaders — you only find them by reading content.

4.5 File-level vs project-level (Q3)




Figure 4: Content type: by-file vs by-repo

	E1 — by-file (n=982)	E2 — by-repo (n=706)
has FEAT marker	80 %	67 %
config-or-dsl	85 %	76 %
other (non-FEAT)	13 %	16 %
GUI-generated	76 %	54 %
hand-edited	6 %	22 %
distinct origins	173	757

Key finding — sampling frame. Sampling **one file per repository** surfaces more diversity: more non-FEAT content and far more **hand-edited** designs (22 % vs 6 %). Sampling **by file** over-weights the auto-generated designs that a few large studies emit in bulk (982 files came from only 173 repos).

4.6 Provenance — a neuroimaging corpus (Q4)

21 800 contents from 757 repositories, led by GitHub with a strong neuroscience signal (`git.mpib-berlin.mpg.de` = Max Planck Institute). The top repositories are pure neuroimaging: **OpenNeuroDatasets** (the canonical fMRI/BIDS dataset hub), **DiedrichsenLab**, **Opioid_Reward_fMRI**, fMRI-task repos (`flashtask`, `PacManStim`).

Key finding — provenance. `.fsf` on SWH is a **real research corpus** from the neuroimaging community. Unlike some extensions, **no synthetic fixture dominates**: the top repository is 18 % of contents (a real study), and the concentration is driven by studies auto-generating one design




Figure 5: Forge provenance




Figure 6: Top repositories

per subject x run.

4.7 Method validation (Q4)

The zero-API reclassifier (label = “does the file contain `set fmri()`”) agrees with the LLM judge on “is this FEAT” for $1\,654 / 1\,662 = 100\%$. A single- token signature is enough to split FEAT from non-FEAT archive-wide at zero cost.

Lesson — bootstrap a cheap classifier from the LLM. Use the (expensive) judge as an oracle to distil a deterministic rule, then *measure* its agreement. For `.fsf` the `set fmri()` marker is essentially perfect.

5. Key findings & lessons (summary)

Findings. 1. `.fsf` ~ **FSL FEAT** fMRI design files — a neuroimaging **configuration** (99 % not a general-purpose program) **expressed in Tcl**. 2. The extension is **polysemous**: also GLSL fragment shaders, plus a large git-annex/DataLad **pointer** artefact. 3. FEAT files are mostly **first-level, task-GLM, GUI-generated**. 4. **By-repo** sampling reveals more hand-edited, more diverse content than by-file. 5. Provenance is a **real neuroimaging corpus** (OpenNeuro, labs, FMRI), no synthetic fixture.

Lessons. 1. Ask “*what is this content and what does it relate to?*” — not “is it language X” — because extensions are polysemous. 2. A `content_type` axis + a **related_languages** list is what exposes both the dominant meaning and the tail. 3. Contamination is **ecosystem-specific**; only content inspection finds it. 4. Report **file-level and project-level** frames separately.

6. Toward a general methodology

This study of `.fsf` is one instance of a general question — *what is actually in file-extension X on Software Heritage?* — and a reusable pipeline for answering it (sample -> fetch -> deterministic indicators -> enum LLM-judge -> cheap reclassifier -> provenance -> human review). We deliberately applied the *same* method to a very different extension, `.cbl/.CBL` (**COBOL**, `docs/cobol_swh_study.md`), to check that it works symmetrically:

	<code>.fsf</code> (this study)	<code>.cbl/.CBL</code> (COBOL)
verdict	a non-PL configuration (FSL FEAT / Tcl)	a real programming language
purity	mostly clean; polysemy = GLSL shaders + git-annex pointers	~46 % contamination (synthetic fixture + comic-book lists)
cheap classifier vs judge	100 % (is-FEAT)	F1 0.88 (is-COBOL)
provenance	real neuroimaging corpus, no fixture	one synthetic fixture = 40 % of contents
deliverables	report + figures + review app + rules	idem

Lesson — the method generalises. The same pipeline characterises both a programming language *and* a domain configuration format, distinguishes genuine content from contamination in each, and surfaces provenance — evidence that “characterise extension X on SWH” can be a general, reproducible methodology rather than a one-off.

7. Limitations

- **git-annex pointers** inflate the FEAT-looking share at the *file* level and under-represent real design content (the actual files are annexed off-SWH).

- The ~21 non-FEAT programs are a small window on the shader/other tail; a targeted sweep of `.fsf` files *without* `set_fmri()` would size it precisely.
- Judge accuracy is validated against the deterministic marker (100 % on is-FEAT) but not yet against human ground truth — the fsf review app (`tools/fsf/review_app.py`, with group-assertion rules) exists for that.
- 757 origins is the whole by-repo ceiling for this extension.

8. Reproducibility & artefacts

```
python3 -m tools.fsf.study --sample --n 1000 --seed 5
SWH_TOKEN=... OPENROUTER_API_KEY=... python3 -m tools.fsf.study --run --judge
python3 -m tools.fsf.study --report
python3 -m tools.fsf.make_figures
python3 -m tools.fsf.review_app                # http://127.0.0.1:8767
```

Artefacts: `data/derived/fsf_study/` (`worklist_all.csv`, `reports/<sha>.json`, `summary_fsf.json`), `fsf_files+origin.csv` (graph provenance), figures under `docs/assets/fsf/`. Toolkit: `tools/fsf/`.

Model: claude-sonnet-4.6 · 1 662 judged · ~\$31 · study authored with Claude Code.